

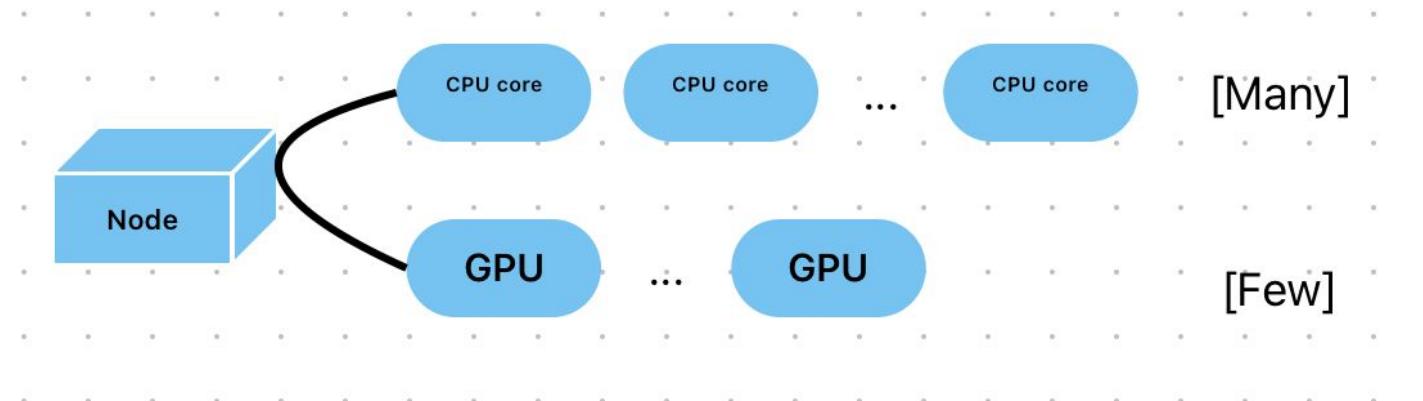
# Slurm tutorial

Bohan Jiang

# What is cluster?

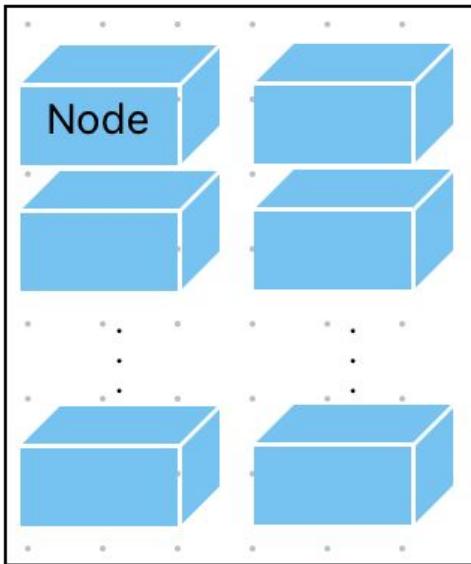
- A cluster is a group (or groups) of computers interconnected to each other, and perform computations for requested tasks.
- Hierarchy
  - a. Cluster
  - b. Partitions
  - c. Nodes

# Node



# Partition

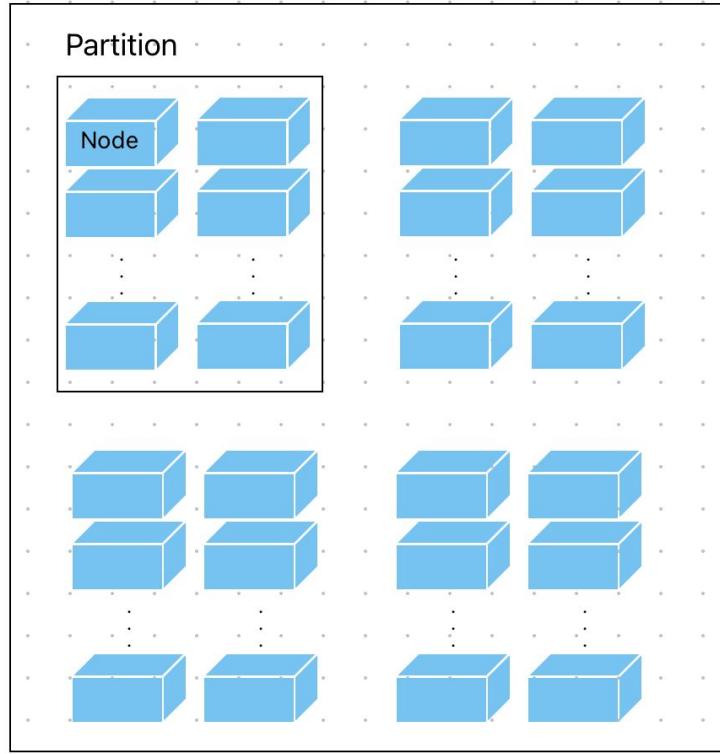
## Partition



- Partition is a group of connected nodes that serves a specific purpose
  - Interactive partition for debugging
  - GPU partition training deep neural models

# Cluster

Cluster



- So, a cluster is a group of partitions
- Partitions may or may not communicate with each other

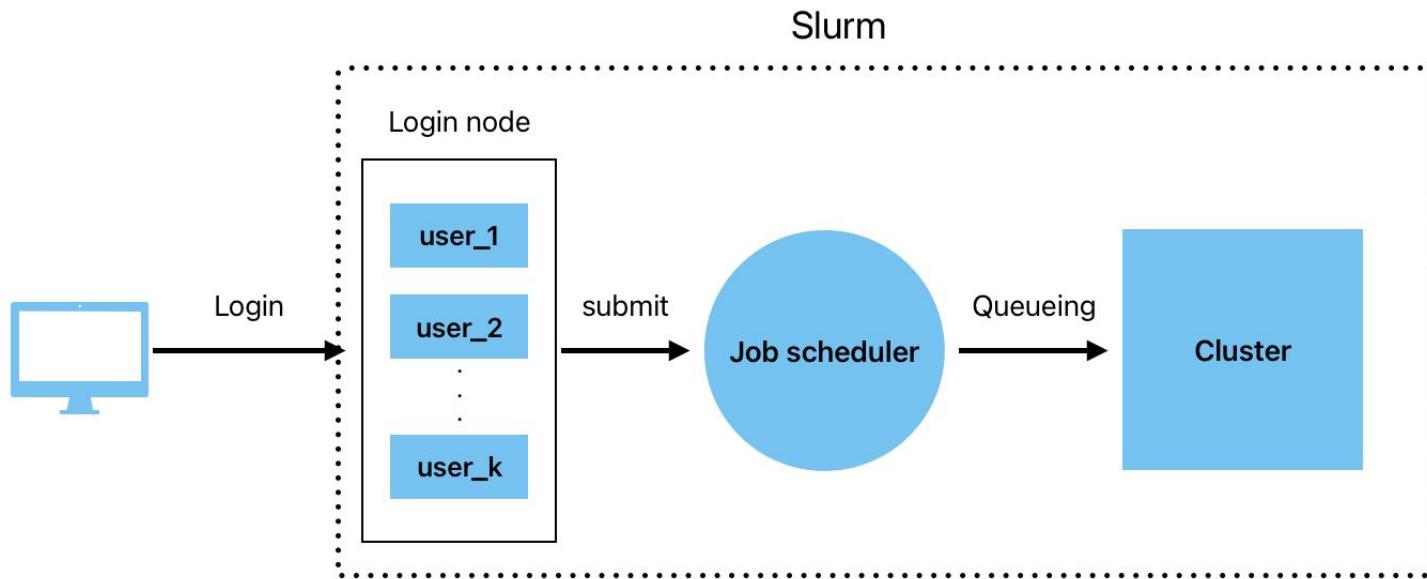
# Brandeis CS student gpu cluster (SGC)

- 2 partitions
  - regular
    - student-gpu-[ 001-004 ]
  - gpu48g
    - student-gpu-[ 003-004 ]

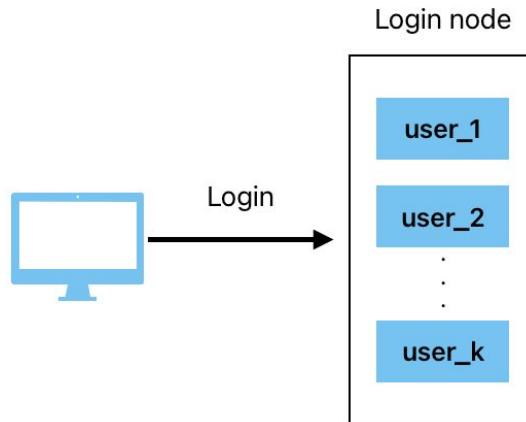
# Slurm

- Slurm is a queuing system that effectively manages the cluster to process submitted job tasks
- That means,
  - You don't have to monitor system usage before running your job
  - On the other hand, you're not the primary user in this world, so your request may be delayed or rejected if it's too resource demanding

# Overview



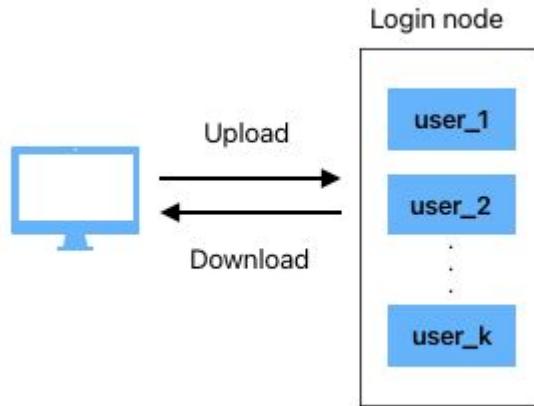
# Login node



- Slurm provides a shared node for all users to log in from their local machines via ssh
  - On campus: direct ssh
  - Off campus: you need VPN
- This shared node is called “login node”, it’s used for submitting job scripts
- Login node has limited resource, so each user has even more limited resource (~2G memory), so DO NOT run any computation in this node

*(For VScode users, the “remote-ssh” extension is not working here, as it would install softwares in login node which takes more memory than assigned)*

# How to edit and run your code?

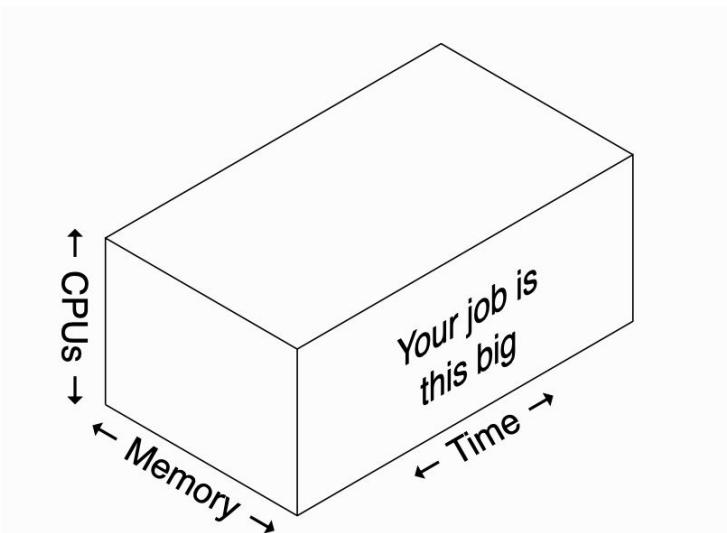


- Solutions
  1. Edit in your local, and run in the cluster

OR

  2. Edit and run in your local in a small scale, and scale it up in the cluster
- Tools
  - scp command
  - SFTP in IDE (e.g. pycharm)

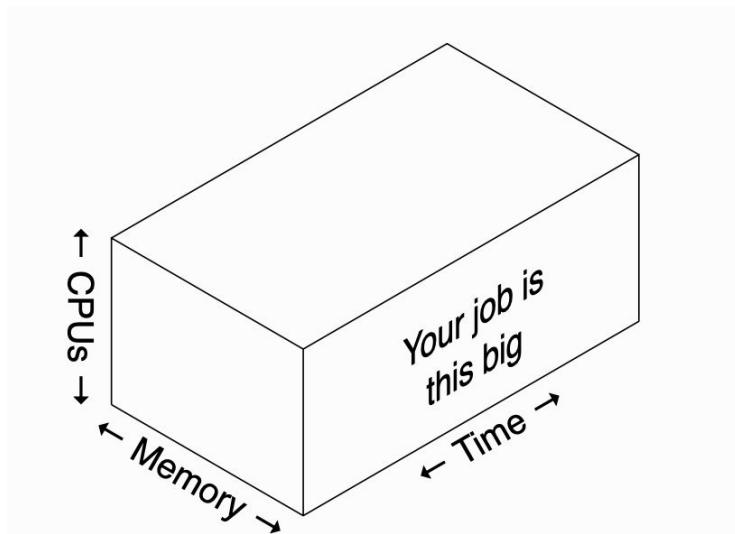
# Submit a job



- 2 scripts:
  - Your root/main script to automate your tasks
  - A “special” bash script
- 1 command:
  - `sbatch <bash_script>`

Source: <https://scicomp.aalto.fi/triton/tut/slurm/>

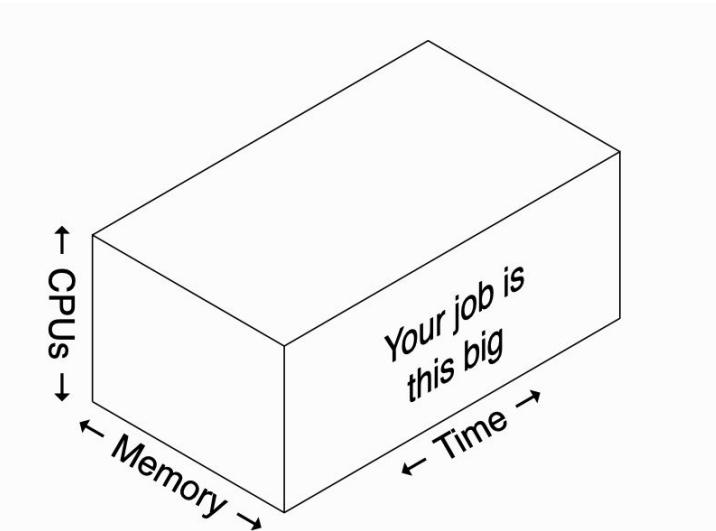
# Submit a job



- Before submission, consider sources you need:
  - Memory
  - Number of CPU cores
  - Number of GPUs
  - How long your job would take

Source: <https://scicomp.aalto.fi/triton/tut/slurm/>

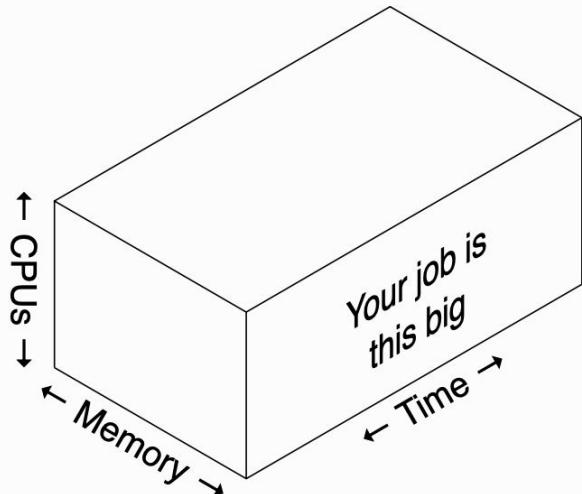
# Submit a job



- Configure the “special” bash script
- Assume you request
  - 4 CPU cores
  - 8G memory
  - 1 GPU device
  - Expect to run 20 mins

Source: <https://scicomp.aalto.fi/triton/tut/slurm/>

# Submit a job



```
#!/bin/bash

#SBATCH --job-name demo
#SBATCH --ntasks 1
#SBATCH --cpus-per-task 4
#SBATCH --mem 8G
#SBATCH --gres 1
#SBATCH --time 00:20:00

cd <WORK_DIR>      # (optional)
export <ENV_VAR>    # (optional)
python main.py
```

Source: <https://scicomp.aalto.fi/triton/tut/slurm/>

# More commands

- `squeue`: show information about job(s) in the queue
  - `squeue -u`: only show info about your job(s) in the queue
- `scancel`: cancel a job
  - `scancel <job_id>`: cancel a job by specifying its id
- `sacct`: show information about your current and previous jobs
- `sinfo`: show availability for nodes in the cluster

# Debug results

- Pop quiz: how do you quickly know if your submitted job is done?
- Your job is done because of 2 reasons:
  - It's executed successfully
  - It's terminated due to errors
- No interactive session once you've submitted your job, where do you find stderr/stdout message?

```
#SBATCH -o <stdout_path>
#SBATCH -e <stderr_path>
```

- These files are often called log files, python has logging module which is more powerful than print statement
- Open them using a text editor or less/cat command

# Demo: useful commands

# Demo: CPU job

# Demo: GPU job